



Docker Deployment Document



1. Objective	3
1.1 About Rasa Chatbot	3
2. Prerequisites.....	3
3. How to get started	4
3.1 Clone the Repository.....	4
3.2 Actions after cloning the repository	6
4. Verification/Validation.....	7
4.1 Validate using Post API.....	7
4.2 Validate using curl on console	7
4.2 Verification from running containers.....	8
5. Container actions.....	9
5.1 Inside container	9
5.2 chat history inside container	9
5.3 stopping the containers	10
6. Running docker process in background	11
7. Rasa Training	11
8. Re-Deploying in the case of new commits	11

1. Objective

This document describes the Deployment Details of AI-powered Rasa Chatbot. It focuses on building the necessary Docker images, composing the network, and creating containers. This document is meant for the Software Developer and Maintenance Team members.

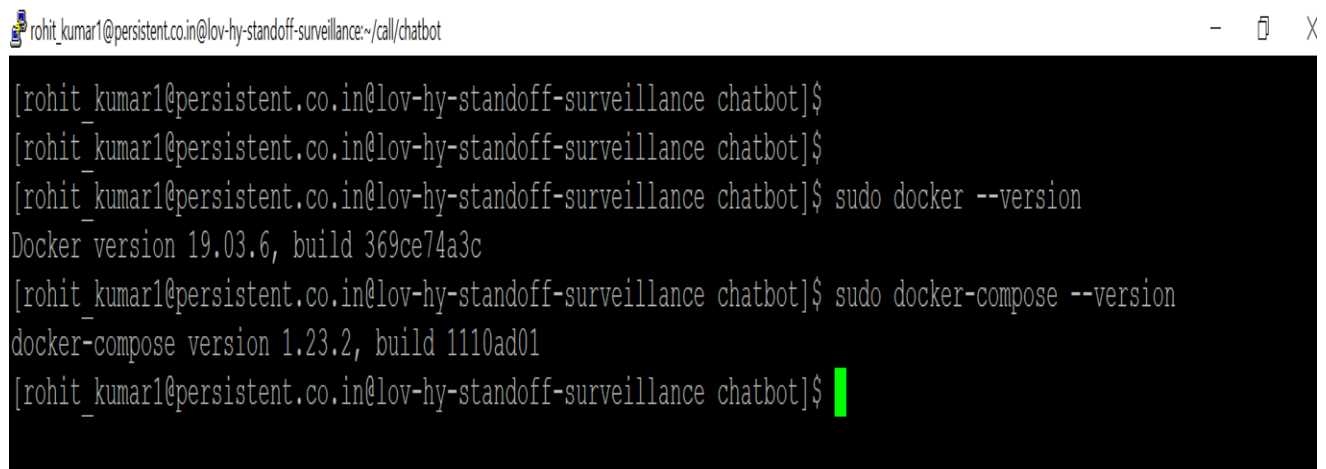
1.1 About Rasa Chatbot

The AI-Powered Rasa Chatbot makes use of two services/servers, specifically Rasa Server and the Action Server. Docker Compose allows us to run both of the servers simultaneously while interacting with each other, maintain the chat history in the container, expose the host over API, and many other operations.

2. Prerequisites

Docker (version 18 or above) and Docker Compose must be installed on the Server/VM.

- To check Docker version: **docker --version**
- To check docker-compose version: **docker-compose --version**



```
rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance:~/call/chatbot
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance chatbot]$
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance chatbot]$
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance chatbot]$ sudo docker --version
Docker version 19.03.6, build 369ce74a3c
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance chatbot]$ sudo docker-compose --version
docker-compose version 1.23.2, build 1110ad01
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance chatbot]$
```

3. How to get started

This section covers the steps needed to create the Compose network, set up both servers, and set up the intercommunication between them.

3.1 Clone the Repository

Clone the repository from the GitHub: **git clone** <https://github.com/NSSAC/chatbot.git> .In order to clone you must have the access to the NSSAC/chatbot repository.

```
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance awe]$ git clone https://github.com/NSSAC/chatbot.git
Cloning into 'chatbot'...
Username for 'https://github.com': rohit-persistent
Password for 'https://rohit-persistent@github.com':
remote: Enumerating objects: 2095, done.
remote: Counting objects: 100% (2095/2095), done.
remote: Compressing objects: 100% (1763/1763), done.
remote: Total 2095 (delta 337), reused 2072 (delta 324), pack-reused 0
Receiving objects: 100% (2095/2095), 43.68 MiB | 9.59 MiB/s, done.
Resolving deltas: 100% (337/337), done.
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance awe]$ cd chatbot/
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance chatbot]$
```

The cloned repository will have Dockerfile and docker-compose.yml file which contains all the rasa image, rasa-dependencies, network and volumes information. Below is the snapshot of the Dockerfile:

```
rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance:~/git_rasa
Base image
FROM ubuntu

#Extend the official Rasa SDK image
FROM rasa/rasa-sdk:1.10.2

MAINTAINER "Rohit Kumar"

#Change back to root user to install dependencies
USER root

#Work Directory
WORKDIR /app

#Installing git and cloning the data from the gitHub
RUN apt-get update && apt-get upgrade -y && apt-get install -y git
RUN git clone https://2b863d8ba6c117a331f107e08241d33a95c572a2:x-oauth-basic@github.com/NSSAC/chatbot.git

#Adding the codebase to our work directory
RUN cp -R ./chatbot/* ./


# Install extra requirements for actions code, if necessary (uncomment next line)
RUN pip install -U pip && pip install word2number && pip install dateparser
RUN pip3 install pandas

# By best practices, don't run the code with root user
USER 1001

CMD ["start", "--actions", "actions", "--debug"]
~
"Dockerfile" 29L, 801C 1,1 All
```

docker-compose.yml file is responsible for the internal port mappings, creating a compose network for both action and rasa server, and the volumes information.

Below is the snapshot of the docker-compose.yml file:

 rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance:~/awe/chatbot

```
version: '3.0'
services:
  rasa:
    image: rasa/rasa:1.10.11-full
    ports:
      - 5005:5005
    volumes:
      - ./:/app
    command:
      - run
      - --enable-api
      - --endpoints
      - endpoints.yml
      - --cors
      - "*"

  app:
    image: rasa/rasa-action-server
    volumes:
      - ./actions:/app/actions
    ports:
      - 5055:5055
```



~
~
~

"docker-compose.yml" 23L, 358C



3.2 Actions after cloning the repository

- Go to the cloned directory 'chatbot' where Dockerfile and docker-compose.yml file exists
- Run **docker build -t rasa/rasa-action-server .** at the level where the Dockerfile exists. It will create a custom rasa action server image. The name of the image must be **rasa/rasa-action-server** or, if you want a custom image name, you also need to change the image name in the docker-compose.yml file. You can see the images by running the **docker images** command.
- As in the snapshot below, rasa/rasa-action-server is our image.

```
Successfully tagged rasa/rasa-action-server:latest
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance chatbot]$ sudo docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
rasa/rasa-action-server	latest	ec7eb136852a	About a minute ago	707MB
python	latest	28a4c88cdbbf	2 weeks ago	882MB
rasa/rasa	1.10.11-full	e3da25b13223	4 weeks ago	2.72GB
ubuntu	latest	4e2eef94cd6b	5 weeks ago	73.9MB
rasa/rasa-sdk	1.10.2	97eb780f0ce2	3 months ago	196MB

- Now, run **docker-compose up** to start both rasa and action servers. This command will create an internal network, start both the servers simultaneously and interconnect them and you will be able to see the logs on the console. In order to run rasa in the background refer section 6.

```
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance chatbot]$ sudo docker-compose up
Creating network "chatbot_default" with the default driver
Creating chatbot_rasa_1 ... done
Creating chatbot_app_1 ... done
Attaching to chatbot_app_1, chatbot_rasa_1
app_1 | 2020-09-25 10:49:21 INFO      rasa_sdk.endpoint - Starting action endpoint server...
app_1 | 2020-09-25 10:49:21 INFO      rasa_sdk.executor - Registered function for 'action count'.
app_1 | 2020-09-25 10:49:21 INFO      rasa_sdk.endpoint - Action endpoint is up and running on http://localhost:5055
app_1 | 2020-09-25 10:49:21 DEBUG      rasa_sdk.utils - Using the default number of Sanic workers (1).
rasa_1 | 2020-09-25 10:49:38 INFO      root - Starting Rasa server on http://localhost:5005
rasa_1 | 2020-09-25 10:50:08 INFO      rasa.nlu.components - Added 'SpacyNLP' to component cache. Key 'SpacyNLP-en_core_web_md'.
rasa_1 | 2020-09-25 10:50:08.958140: E tensorflow/stream_executor/cuda/cuda_driver.cc:351] failed call to cuInit: UNKNOWN ERROR (303)
rasa_1 | /opt/venv/lib/python3.7/site-packages/sklearn/externals/joblib/_init_.py:15: FutureWarning: sklearn.externals.joblib is deprecated in 0.21 and will be removed in 0.23. Please import this functionality directly from joblib, which can be installed with: pip install joblib. If this warning is raised when loading pickled models, you may need to re-serialize those models with scikit-learn 0.21+.
rasa_1 | warnings.warn(msg, category=FutureWarning)
rasa_1 | 2020-09-25 10:50:23 INFO      root - Rasa server is up and running.
```

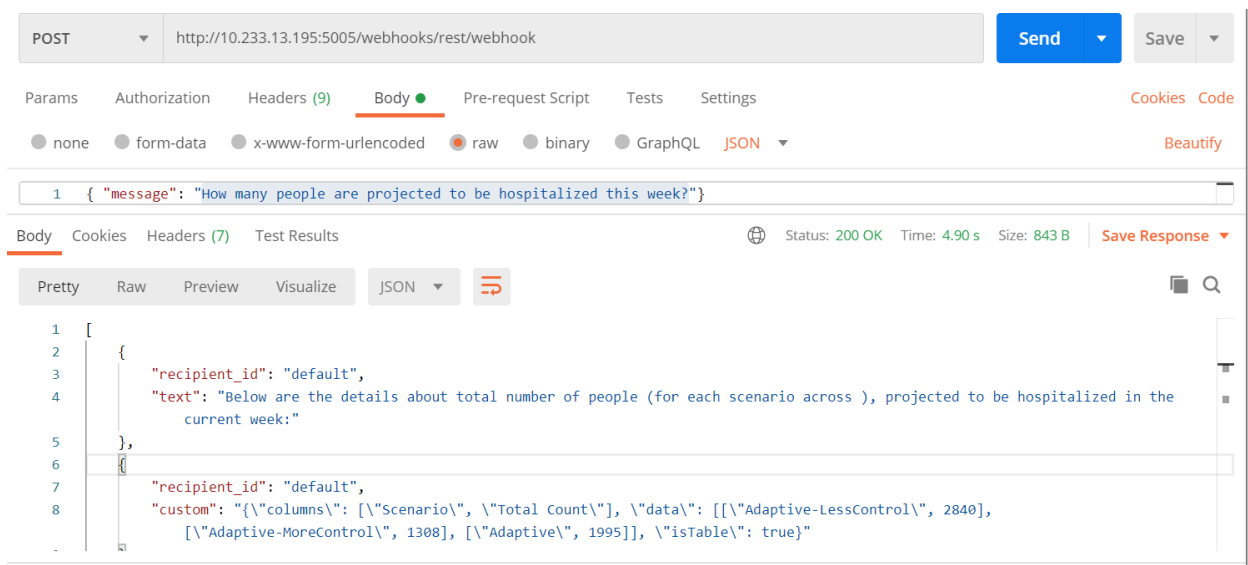
4. Verification/Validation

4.1 Validate using Post API

Now that both of the servers are up and running, we can access the docker service by making a POST API request to the docker container using the details below:

- <http://<host-ip>:5005/webhooks/rest/webhook>
- Content-type: application/json
- Body example:

```
{ "message": "How many people are projected to be hospitalized this week?" }
```



4.2 Validate using curl on console

We can also verify by sending the messages to the bot from the console as shown below:

```
curl --request POST \  
  
  --url http://localhost:5005/webhooks/rest/webhook \  
  
  --header 'content-type: application/json' \  
  
  --data '{  
  
    "message": "How many people are projected to be hospitalized this week?"  
  
  }'
```

Your chatbot should answer similar to this:

```
[{"recipient_id":"default","text":"Below are the details about total number of people  
(for each scenario across ), projected to be hospitalized in the current week:"},  
  
{"recipient_id":"default","custom":{"columns\":[\"Scenario\", \"Total Count\"],  
\"data\":[[\"Adaptive-LessControl\", 2840], [\"Adaptive-MoreControl\", 1308],  
[\"Adaptive\", 1995]], \"isTable\": true}},  
  
{"recipient_id":"default","text":"What else do you want to check about?",  
  
"buttons":[{"payload":"\\/hospitalization","title":"Hospitalization"},  
  
{"payload":"\\/occupied_beds","title":"Occupied  
Beds"}, {"payload":"\\/about_mrdd","title":"About MRDD"}]}
```

4.2 Verification from running containers

To list out all of the running containers, run the **docker ps -a** command. In the snapshot below, both the rasa-action-server and rasa server containers can be seen up and running

```
[rohit_kumar@persistent.co.in@lov-hy-standoff-surveillance files]$  
[rohit_kumar@persistent.co.in@lov-hy-standoff-surveillance files]$  
[rohit_kumar@persistent.co.in@lov-hy-standoff-surveillance files]$  
[rohit_kumar@persistent.co.in@lov-hy-standoff-surveillance files]$ sudo docker ps -a  
CONTAINER ID        IMAGE               COMMAND             CREATED             STATUS              PORTS              NAMES  
c5a3440d93ea        rasa/rasa-action-server  "/entrypoint.sh sta..." About a minute ago  Up About a minute  0.0.0.0:5055->5055/tcp  chatbot_app_1  
676a163140ee        rasa/rasa:1.10.11-full  "rasa run --enable-a..." About a minute ago  Up About a minute  0.0.0.0:5005->5005/tcp  chatbot_rasa_1  
[rohit_kumar@persistent.co.in@lov-hy-standoff-surveillance files]$
```


5.1 Inside container

```
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance files]$ sudo docker exec -it c5a3440d93ea bash
I have no name!@c5a3440d93ea:/app$ ls
CHANGELOG.rst      Makefile           changelog          data                entrypoint.sh      poetry.lock        rasa_sdk.egg-info
CustomEntityExtractor.py  README.md          chatbot            docker-compose.yml  examples            pyproject.toml     scripts
Dockerfile          __init__.py        config.yml          domain.yml           listfile.txt        rasa_core_sdk      setup.cfg
LICENSE.txt          actions            credentials.yml     endpoints.yml        models              rasa_sdk            tests
I have no name!@c5a3440d93ea:/app$
```

Chat history with chatbot is maintained in the directory named logs inside in the rasa action container, and to see the contents of the chat history we need to go to that directory and see as shown in the snapshot below.

```
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance ~]$
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance ~]$ sudo docker ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
4b46e4c89dbb	rasa/rasa-action-server	"./entrypoint.sh sta..."	3 days ago	Up 3 days	0.0.0.0:5055->5055/tcp	chatbot_app_1
a0423079a52a	rasa/rasa:1.10.11-full	"rasa run --enable-a..."	3 days ago	Up 3 days	0.0.0.0:5005->5005/tcp	chatbot_rasa_1

```
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance ~]$ sudo docker exec -it 4b46e4c89dbb bash
I have no name!@4b46e4c89dbb:/app$ ls
```

CHANGELOG.rst	Makefile	chatbot	docker-compose.yml	examples	poetry.lock	rasa_sdk.egg-info
'Docker Deployment Document.docx'	README.md	config.yml	domain.yml	log_file_parsing_chatbot.py	pyproject.toml	scripts
Dockerfile	actions	credentials.yml	endpoints.yml	logs	rasa_core_sdk	setup.cfg
LICENSE.txt	changelog	data	entrypoint.sh	models	rasa_sdk	tests

```
I have no name!@4b46e4c89dbb:/app$ cd logs/
I have no name!@4b46e4c89dbb:/app/logs$ ls
```

README.txt	default_log.txt
------------	-----------------

```
I have no name!@4b46e4c89dbb:/app/logs$ cat default_log.txt
```


6. Running docker process in background

To run the docker compose network in the background run **docker-compose up -d**

```
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance thursdaycall]$  
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance thursdaycall]$ sudo docker-compose down  
Stopping thursdaycall_rasa_1 ... done  
Stopping thursdaycall_app_1 ... done  
Removing thursdaycall_rasa_1 ... done  
Removing thursdaycall_app_1 ... done  
Removing network thursdaycall_default  
[rohit_kumar1@persistent.co.in@lov-hy-standoff-surveillance thursdaycall]$ sudo docker ps -a
```

- To stop this demon process or running rasa in the background run **docker-compose down**
- As it is running in the background, logs of the rasa servers will not appear on the console. If you still want to see the logs of the containers run **docker container logs <container-id>**

7. Rasa Training

To train the Rasa for the newly added intents, or if you don't have any model available to you, then run

docker run --user 1000 -it -v \$(pwd):/app rasa/rasa:1.10.11-full train

8. Re-Deploying in the case of new commits

- If there is a new commit GitHub, then follow the instructions below:
 - Stop the currently running compose network: **ctrl+c**
 - Delete the old rasa-action-server image: **docker rmi rasa/rasa-action-server**
 - Run **docker build -t rasa/rasa-action-server .**
 - Then run **docker-compose up**
- **docker system prune** command removes all the dangling images and stopped containers.
- To inspect or learn more information about a container like ipaddress,port,volume etc.
run docker inspect <container-id>
- Point to remember: Always run build and compose commands at the level where the Dockerfile and docker-compose.yml file is present.

